

$$C = A * B$$

Each process gets a
number of rows: n/P

*

Communicate entire B to all
other processes
(from process 0)

=

Calculated by
process P_i

Calculated by
process P_{i+1}

A

B

C

```
double a[NRA][NCA],    /* matrix A to be multiplied */
      b[NCA][NCB],    /* matrix B to be multiplied */
      c[NRA][NCB];    /* result matrix C */
```

...

/* Parallelize the computation of the following matrix-matrix multiplication.

How to partition and distribute the initial matrices, the work, and collecting final results.

SPMD programming style:

```
*/

for (i=0; i<NRA; i++)
{
    for(j=0; j<NCB; j++)
        for (k=0; k<NCA; k++)
            c[i][j] += a[i][k] * b[k][j];
}
```

```
if (myID == 0)
{
    For i=1..(P-1) Send n/P rows of A and B to P(i);
}
else
{
    Receive n/P rows of A and B from P(0);
}
Calculate part of A(i)*B, and send back
```

What?

```
MPI_Send ( void *data, int count, MPI_Datatype type,  
           int dest, int tag, MPI_Comm comm )
```

To whom?

C

Pointers and contents: &, *, (int *) a, (double *) b, etc.

```
double a;  
MPI_Send(&a, 1, ....);
```

```
double *a;  
MPI_Send(a, 1, ....)
```

```
double b[100];  
MPI_Send(&b[0], 100, ....);  
Or  
MPI_Send(b, 100, ...)
```

Function: call by reference and call by value

MPI processes (with local variables), similarity with call by value

MPI Collective Communications

- MPI_Reduce
- MPI_Allreduce
- MPI_Gather, MPI_Gatherv
- MPI_Scatter, MPI_Scatterv

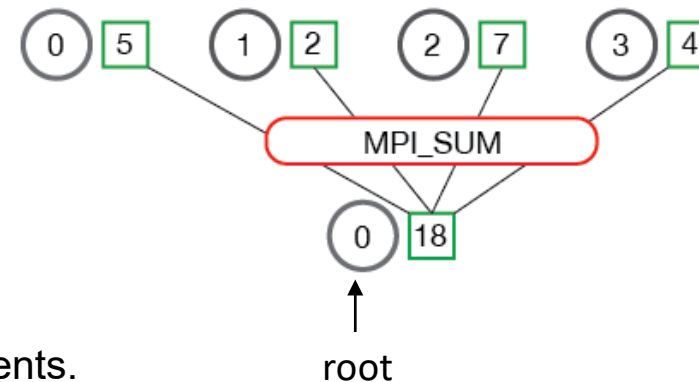
MPI_Reduce

```
MPI_Reduce( void* send_data, void* recv_data, int count,  
MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm  
communicator)
```

The reduction operations defined by MPI include:

- MPI_MAX - Returns the maximum element.
- MPI_MIN - Returns the minimum element.
- MPI_SUM - Sums the elements.
- MPI_PROD - Multiplies all elements.
- MPI_LAND - Performs a logical *and* across the elements.
- MPI_LOR - Performs a logical *or* across the elements.
- MPI_BAND - Performs a bitwise *and* across the bits of the elements.
- MPI_BOR - Performs a bitwise *or* across the bits of the elements.
- MPI_MAXLOC - Returns the maximum value and the rank of the process that owns it
- MPI_MINLOC - Returns the minimum value and the rank of the process that owns it.

MPI_Reduce



Dot product using reduction

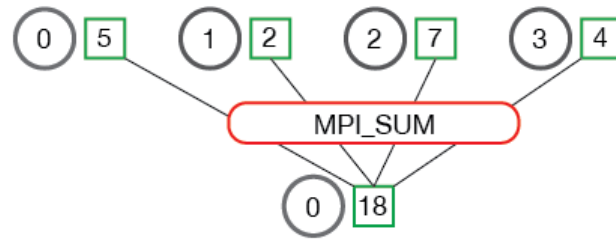
```
int j, m, p, local_m;
float local_dot, dot;
float local_x[100], local_y[100];
MPI_Status status;

MPI_Comm_rank( MPI_COMM_WORLD, &my_rank);
MPI_Comm_size( MPI_COMM_WORLD, &p);
if (my_rank == 0) scanf("%d",&m);
local_m = m/p;
local_dot = 0.0;
for (j=0; j < local_m; j++)
    local_dot = local_dot + local_x[j] * local_y[j];
MPI_Reduce(&local_dot, &dot,1, MPI_FLOAT, MPI_SUM,0, MPI_COMM_WORLD);
```

MPI program for the parallel computation of a scalar product

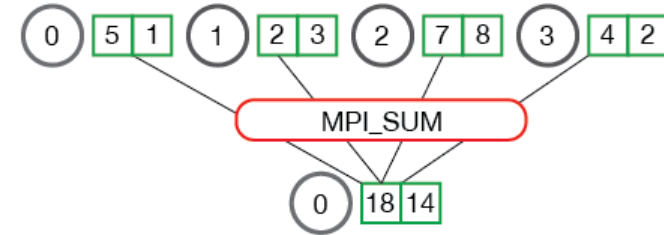
MPI_Reduce and MPI_Allreduce

MPI_Reduce



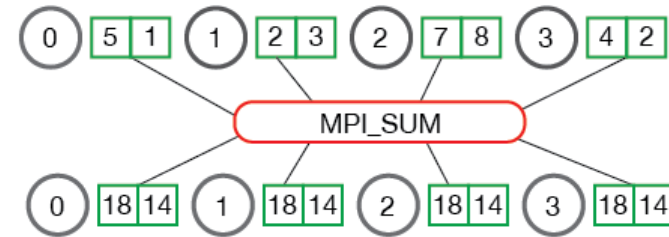
```
MPI_Reduce(&local,&global_sum,1,MPI_INT,  
MPI_SUM,0,MPI_COMM_WORLD)
```

MPI_Reduce



```
MPI_Reduce(&local,&global_sum,2,MPI_INT,  
MPI_SUM,0,MPI_COMM_WORLD)
```

MPI_Allreduce



```
MPI_Allreduce(&local,&global_sum,2,MPI_INT,  
MPI_SUM,MPI_COMM_WORLD)
```

MPI Reduction operation

```
int MPI_Reduce (void *sendbuf,  
               void *recvbuf,  
               int count,  
               MPI_Datatype type,  
               MPI_Op op,  
               int root,  
               MPI_Comm comm)
```

```
double ain[30], aout[30];  
int ind[30];  
struct {double val; int rank;} in[30], out[30];  
int i, my_rank, root=0;  
  
MPI_Comm_rank (MPI_COMM_WORLD, &my_rank);  
for (i=0; i<30; i++) {  
    in[i].val = ain[i];  
    in[i].rank = my_rank;  
}  
MPI_Reduce(in,out,30,MPI_DOUBLE_INT,MPI_MAXLOC,root,MPI_COMM_WORLD);  
if (my_rank == root)  
    for (i=0; i<30; i++) {  
        aout[i] = out[i].val;  
        ind[i] = out[i].rank;  
    }
```

Example for the use
of MPI_Reduce ()
using
MPI_MAXLOC for
the reduction
operator.

MPI_Gather and MPI_Scatter

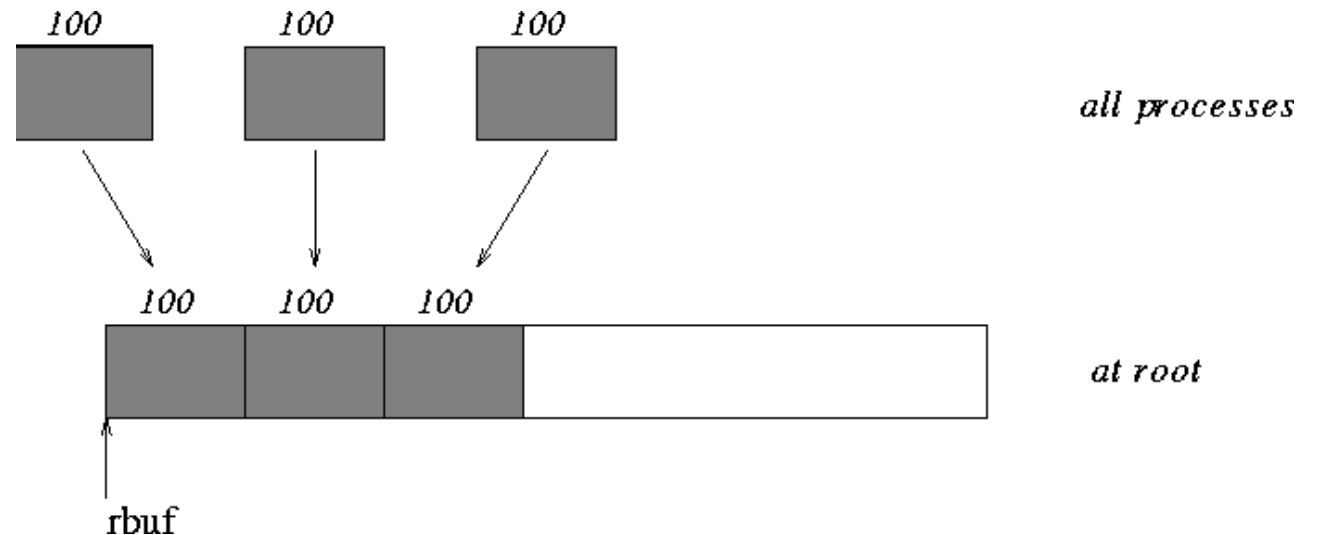
MPI_Gather(const void *sendbuf, int sendcount, MPI_Datatype sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)

Examples from www.mpi-forum.org/docs/mpi-1.1/mpi-11-html/node70.html#Node70

Example:

The root process gathers 100 ints from each process in the group.

```
MPI_Comm comm;
int NrProcs, sendarray[100];
int root, myrank, *rbuf; ...
MPI_Comm_rank( comm, myrank);
if ( myrank == root)
{
    MPI_Comm_size( comm, &NrProcs);
    rbuf = (int *)malloc(NrProcs*100*sizeof(int));
}
MPI_Gather( sendarray, 100, MPI_INT, rbuf, 100, MPI_INT, root, comm);
```

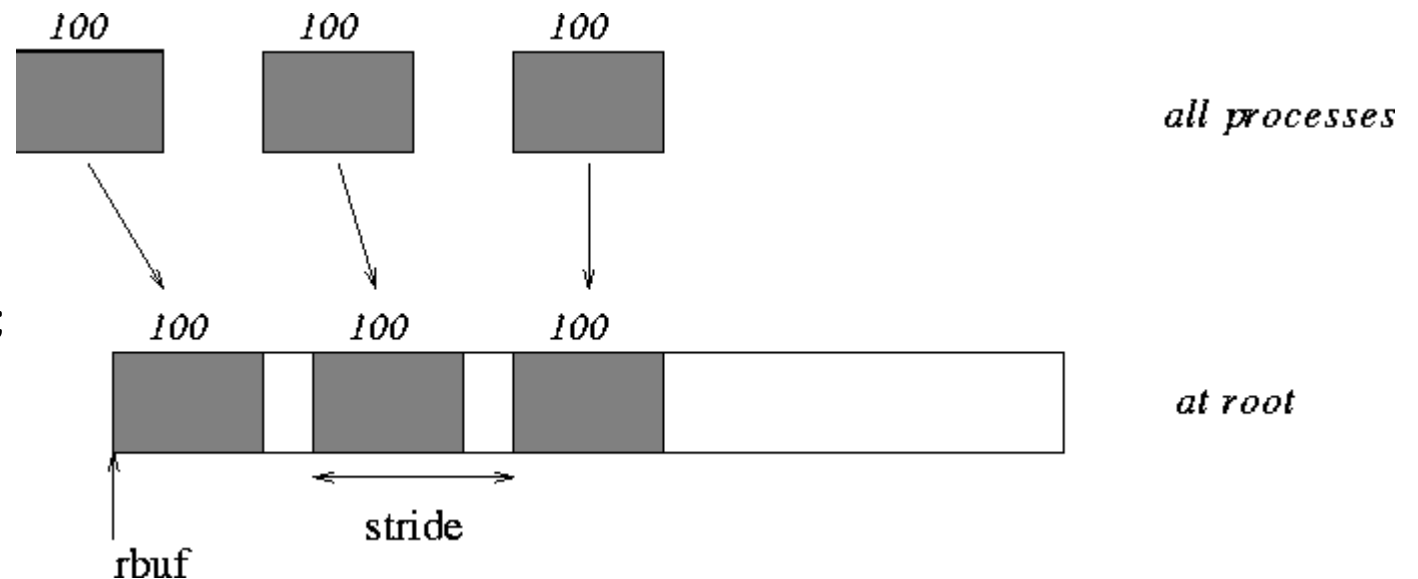


MPI_Gatherv

Example

Now have each process send 100 ints to root, but place each set (of 100) *stride* ints apart at receiving end. Use MPI_GATHERV and the displs argument to achieve this effect. Assume *stride* ≥ 100.

```
MPI_Comm comm;  
int NrProcs, sendarray[100];  
int root, *rbuf, stride;  
int *displs,i,*rcounts;  
...  
MPI_Comm_size( comm, &NrProcs);  
rbuf = (int *)malloc(NrProcs*stride*sizeof(int));  
displs = (int *)malloc(NrProcs*sizeof(int));  
for (i=0; i<NrProcs; ++i)  
{  
    displs[i] = i*stride;  
    rcounts[i] = 100;  
}  
MPI_Gatherv( sendarray, 100, MPI_INT, rbuf, rcounts, displs, MPI_INT, root, comm);
```

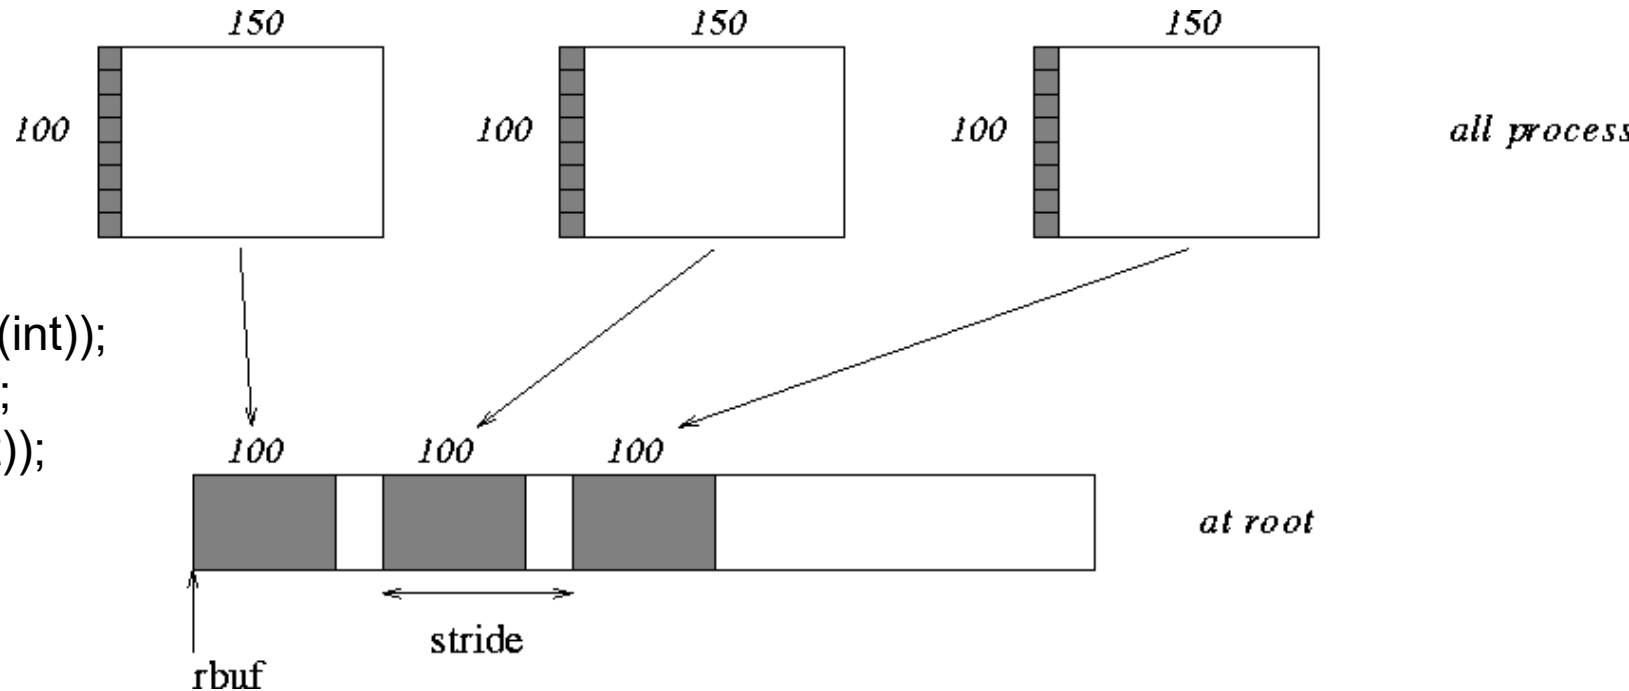


MPI_Gatherv

Example

Same as Example as previous MPI_Gatherv, but now send the 100 ints from the 0th column of a 100×150 int array, in C.

```
MPI_Comm comm;
int NrProcs, sendarray[100][150];
int root, *rbuf, stride;
MPI_Datatype stype;
int *displs, i, *rcounts;
...
rbuf = (int *)malloc(NrProcs*stride*sizeof(int));
displs = (int *)malloc(NrProcs*sizeof(int));
rcounts = (int *)malloc(NrProcs*sizeof(int));
for (i=0; i<NrProcs; ++i) {
    displs[i] = i*stride;
    rcounts[i] = 100;
}
/* Create datatype for 1 column of array */
MPI_Type_vector( 100, 1, 150, MPI_INT, &stype);
MPI_Type_commit( &stype );
MPI_Gatherv( sendarray, 1, stype, rbuf, rcounts, displs, MPI_INT, root, comm);
```



MPI_Scatter

```
MPI_Scatter(void* sendbuf, int sendcount, MPI_Datatype sendtype, void* recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)
```

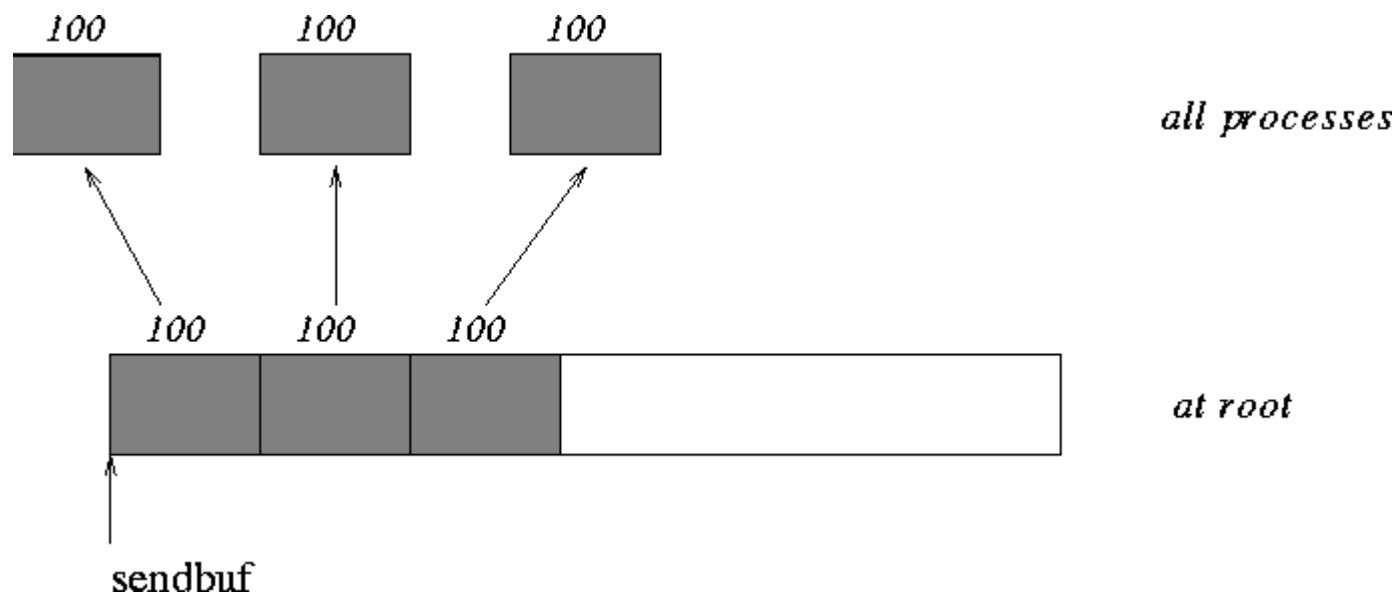
Example

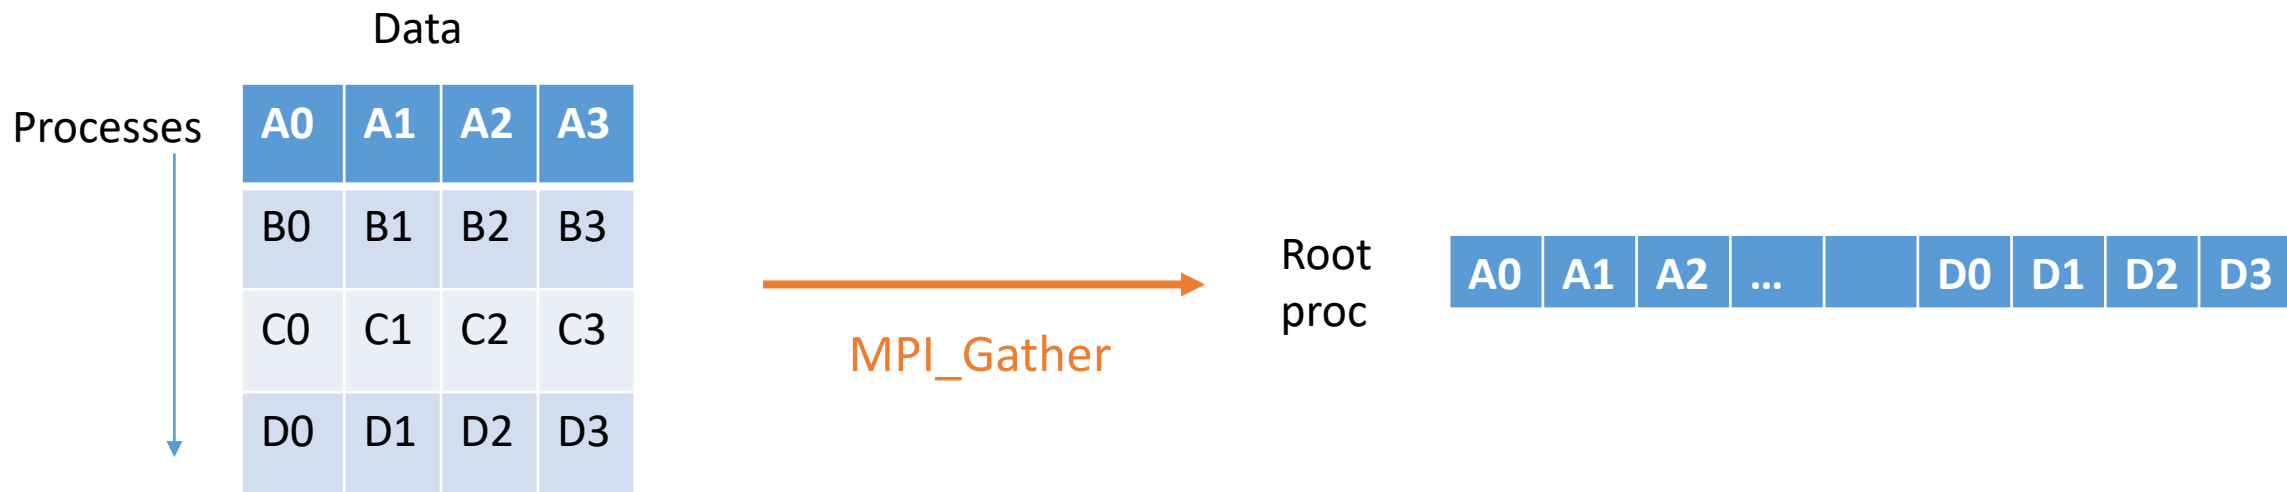
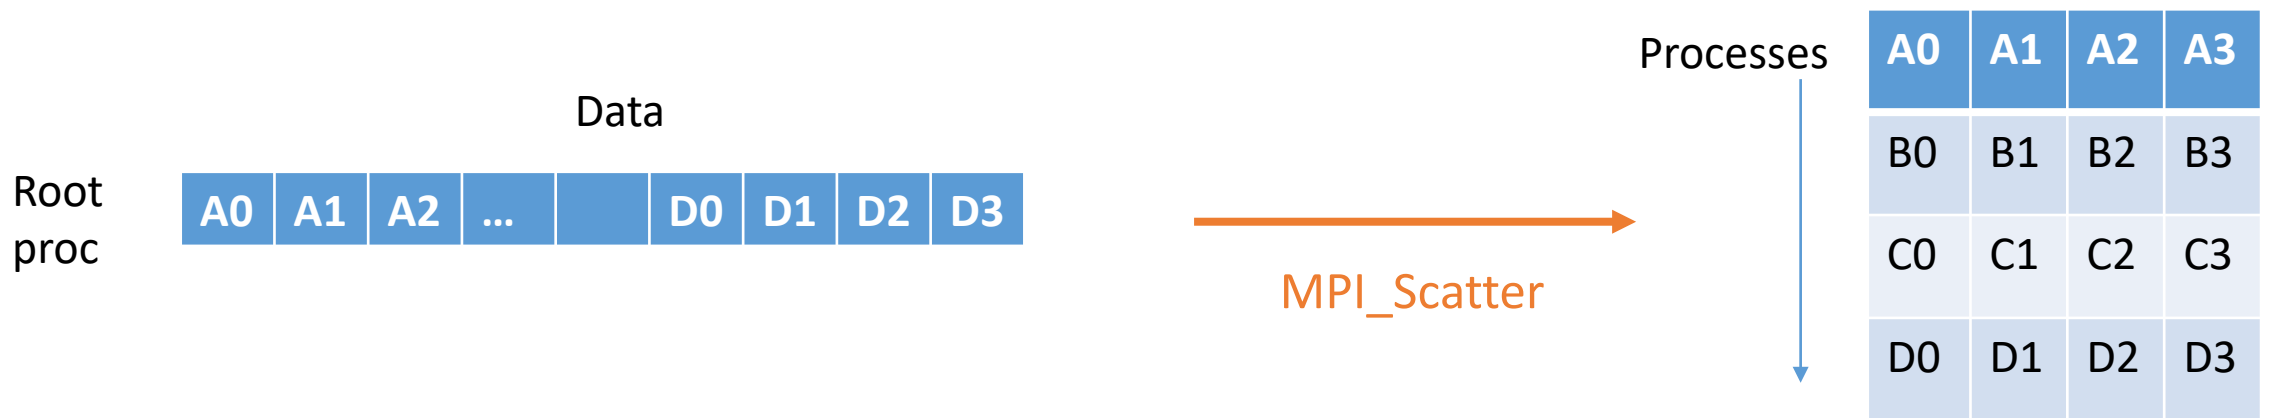
The reverse of Example using [MPI_GATHER](#). Scatter sets of 100 ints from the root to each process in the group.

```
MPI_Comm comm;  
int NrProcs,*sendbuf;  
int root, rbuf[100];  
...
```

```
MPI_Comm_size( comm, &NrProcs);  
sendbuf = (int *)malloc(NrProcs*100*sizeof(int));  
...
```

```
MPI_Scatter( sendbuf, 100, MPI_INT, rbuf, 100, MPI_INT, root, comm);
```



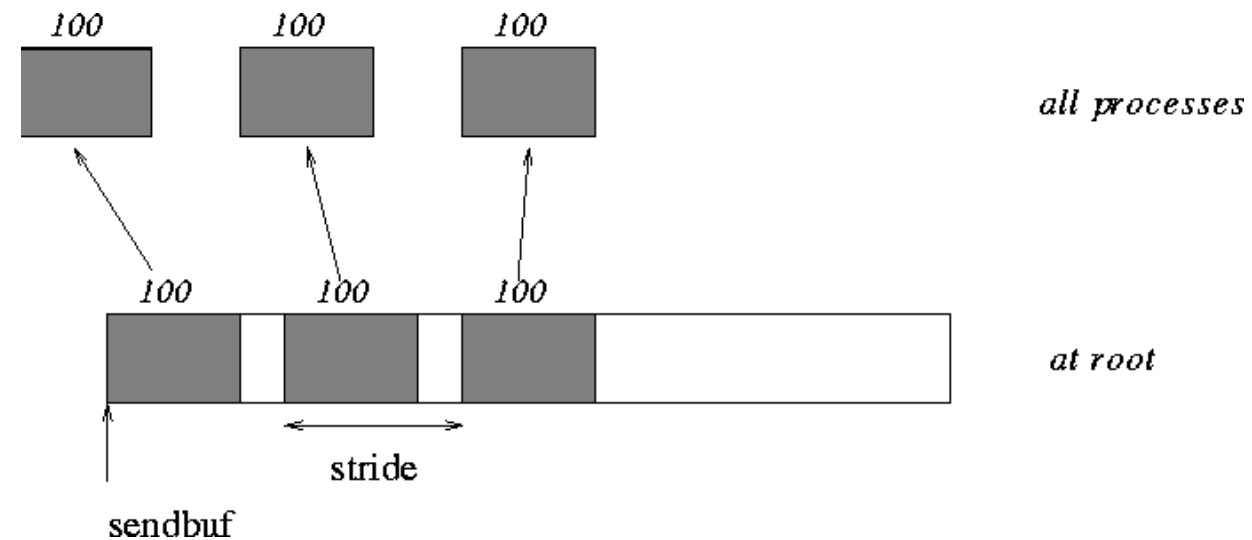


MPI_Scatterv

Example

The reverse of Example using [MPI_GATHERV](#). The root process scatters sets of 100 ints to the other processes, but the sets of 100 are *stride* ints apart in the sending buffer. Assume *stride* ≥ 100.

```
MPI_Comm comm;  
int NrProcs,*sendbuf;  
int root, rbuf[100], i, *displs, *counts;  
...  
MPI_Comm_size( comm, &NrProcs);  
sendbuf = (int *)malloc(NrProcs*stride*sizeof(int));  
...  
displs = (int *)malloc(NrProcs*sizeof(int));  
counts = (int *)malloc(NrProcs*sizeof(int));  
for (i=0; i<NrProcs; ++i) {  
    displs[i] = i*stride;  
    counts[i] = 100;  
}  
MPI_Scatterv( sendbuf, counts, displs, MPI_INT, rbuf, 100, MPI_INT, root, comm);
```



```

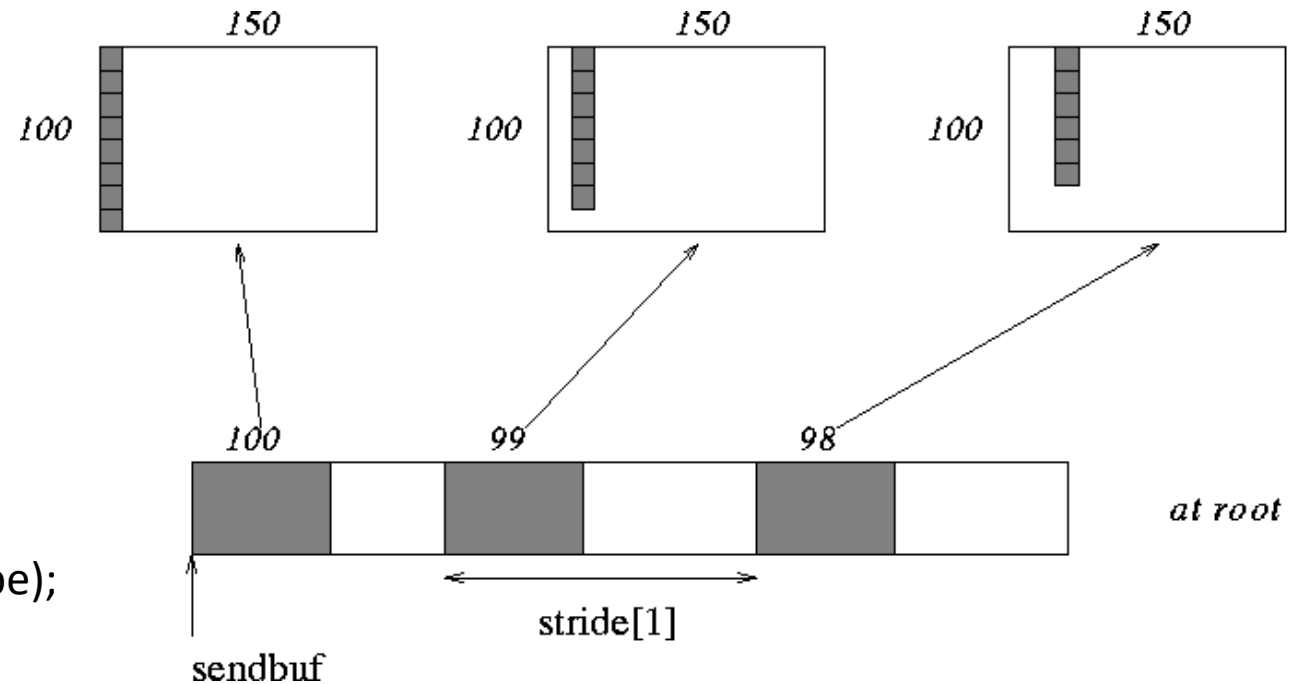
MPI_Comm comm;
int NrProcs, recvarray[100][150], *rptr;
int root, *sendbuf, myrank, bufsize, *stride;
MPI_Datatype rtype;
int i, *displs, *counts, offset;
...
MPI_Comm_rank( comm, &myrank );
stride = (int *)malloc(NrProcs*sizeof(int));
...
/* stride[i] for i = 0 to gsize-1 is set somehow * sendbuf comes from elsewhere */
...
displs = (int *)malloc(NrProcs*sizeof(int));
counts = (int *)malloc(NrProcs*sizeof(int));
offset = 0;
for (i=0; i<NrProcs; ++i) {
    displs[i] = offset;
    offset += stride[i];
    counts[i] = 100 - i;
}
/* Create datatype for the column we are receiving */
MPI_Type_vector( 100-myrank, 1, 150, MPI_INT, &rtype);
MPI_Type_commit( &rtype );
rptr = &recvarray[0][myrank];
MPI_Scatterv( sendbuf, counts, displs, MPI_INT, rptr, 1, rtype, root, comm);

```

MPI_Scatterv

Example

The root scatters blocks of $100-i$ ints into column i of a 100×150 C array. At the sending side, the blocks are $\text{stride}[i]$ ints apart.

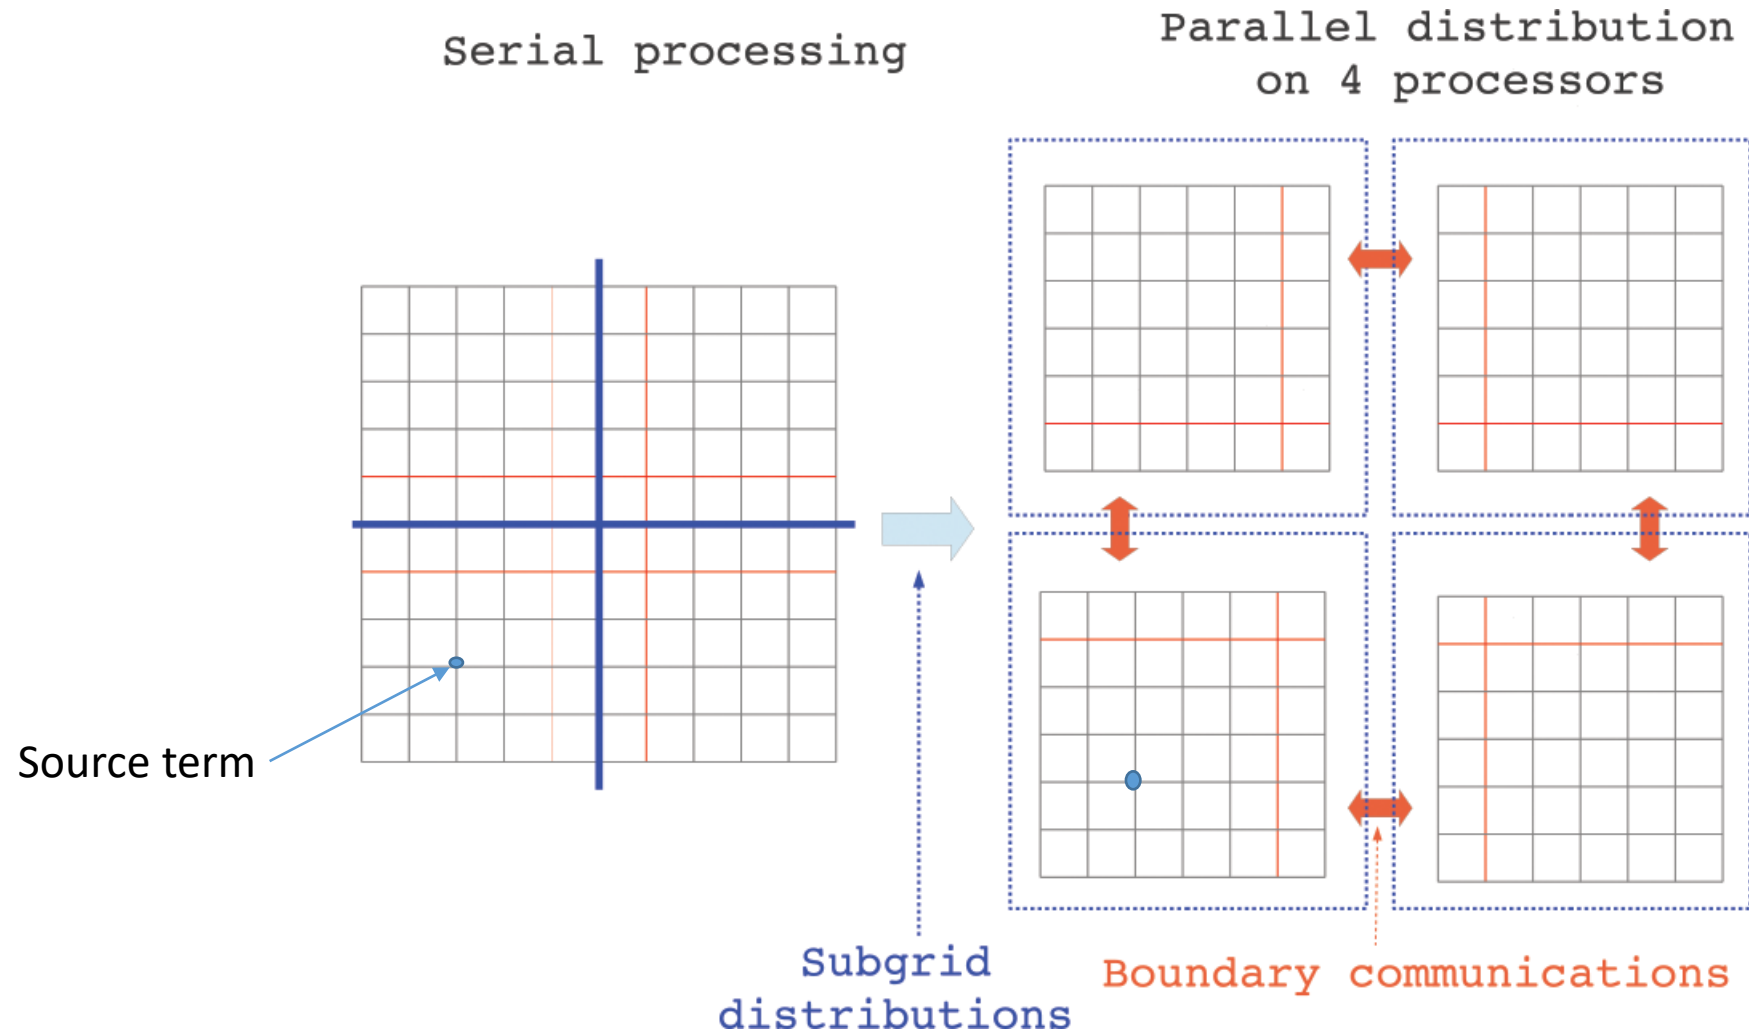


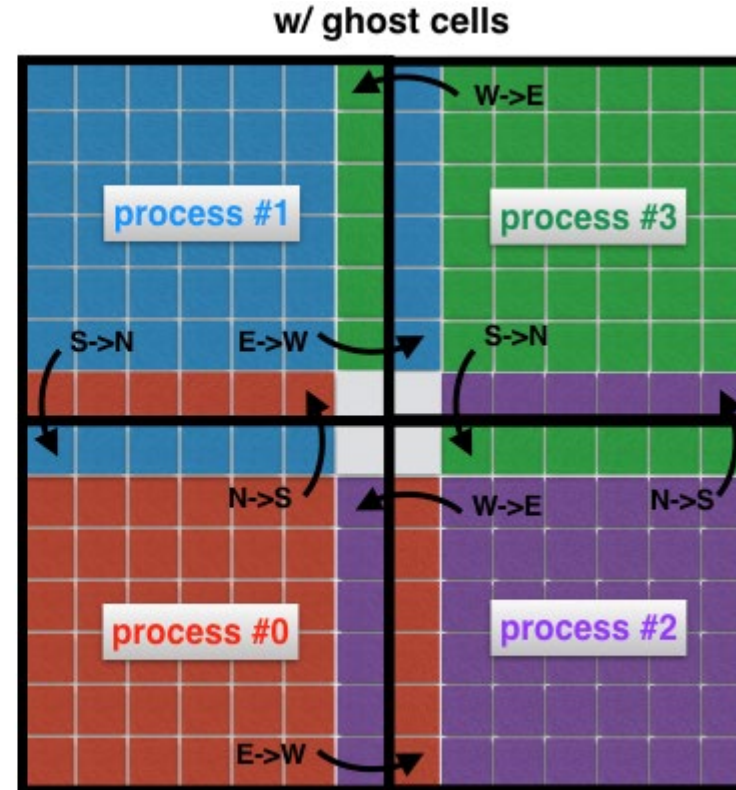
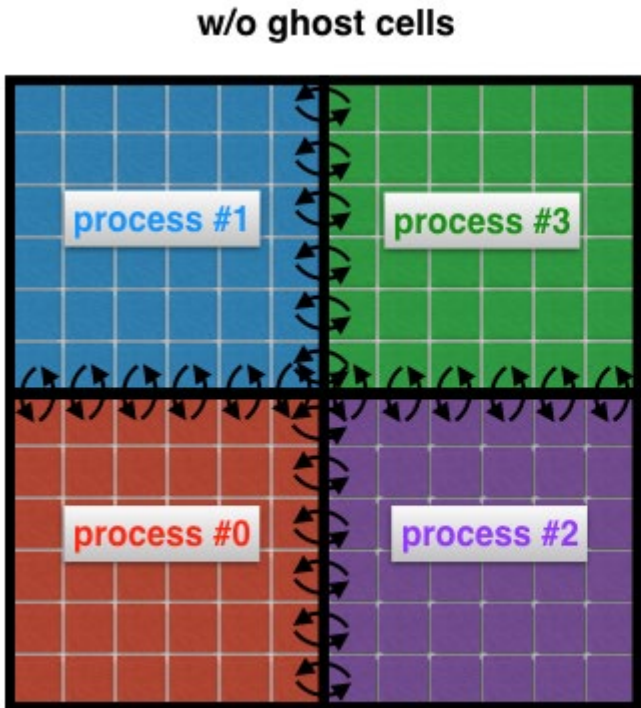
Exercise 1:
Poisson Solver

Ghost cells and HALO-exchange

Process topology (2-D mesh)

Discretization grid and partitioning





Suppose the size of a subgrid is $n_{\text{row}} \times n_{\text{col}}$, the array size increases with ghost cells to $(n_{\text{row}}+2) \times (n_{\text{col}}+2)$.

Row 0 and $n_{\text{row}}+1$ stores the values received from the top and bottom neighbour respectively.

Similar, for column 0 and $n_{\text{col}}+1$...
(All local data start from (0,0)!!)

Communication among processes without (left) and with (right) ghost cells. Without ghost cells, each cell on the boundary of a sub-domain needs to pass its own message to a neighboring process. Using ghost cells allows to minimize the number of messages passed, as many cells belonging to the boundaries of a process are exchanged at once with a single message. Here, for example, Process #0 is passing its entire North boundary to Process #1, and its entire East boundary to Process #2.

2-D Cartesian process topology (for easy neighbor communication)

0 (0,0)	1 (0,1)	2 (0,2)	3 (0,3)
4 (1,0)	5 (1,1)	6 (1,2)	7 (1,3)
8 (2,0)	9 (2,1)	10 (2,2)	11 (2,3)
12 (3,0)	13 (3,1)	14 (3,2)	15 (3,3)

MPI_rank

MPI_Cart_coords

Creating a virtual topology: **comm_2d**

- MPI_COMM_WORLD/oldcomm communicator
- **ndims** dimension of the Cartesian topology
- **dims** integer array (size ndims) that defines the number of processes in each dimension
- **periods** array that defines the periodicity of each dimension – reorder is MPI allowed to renumber the ranks
- **comm_2d** new Cartesian communicator

MPI_Cart_create(MPI_COMM_WORLD, ndims, dims, periods, reorder, &**comm_2d**)

Here, ndims=2, dims[0]=4, dims[1]=4, periods[0]=periods[1]=0, reorder=0

MPI_Comm_rank(comm_2d, **my_rank**, ierror)

MPI_Cart_coords(comm_2d, my_rank, maxdims, **my_coords**, ierror)

Cartesian neighbours

Counting "hops" in the Cartesian grid to allow for e.g., elegant nearest-neighbor communication

`MPI_Cart_shift(comm, direction, displ, source, dest)`

- **comm** Cartesian communicator
- **direction** shift direction (e.g., 0 or 1 in 2D)
- **displ** shift displacement (1 for next cell etc., <0 for "down"/"left" directions)
- **source** rank of source process
- **dest** rank of destination process

With non-periodic grid, *source* or *dest* can land outside of the grid; then `MPI_PROC_NULL` (or sometimes a negative rank number) is returned

0 (0,0)	1 (0,1)	2 (0,2)	3 (0,3)
4 (1,0)	5 (1,1)	6 (1,2)	7 (1,3)
8 (2,0)	9 (2,1)	10 (2,2)	11 (2,3)
12 (3,0)	13 (3,1)	14 (3,2)	15 (3,3)

invisible input argument: `my_rank` of the running process executing:

`MPI_Cart_shift(comm_2d, direction, displ, &prev, &next)`

example on process rank=6	0	+1	5	7
	1	+1	2	10
	0	-1	7	5